




# CUDA C++ BASICS

NVIDIA Corporation



# WHAT IS CUDA?


- ▶ CUDA Architecture
  - ▶ Expose GPU parallelism for general-purpose computing
  - ▶ Expose/Enable performance
- ▶ CUDA C++
  - ▶ Based on industry-standard C++
  - ▶ Set of extensions to enable heterogeneous programming
  - ▶ Straightforward APIs to manage devices, memory etc.
- ▶ This session introduces CUDA C++
  - ▶ Other languages/bindings available: Fortran, Python, Matlab, etc.

# INTRODUCTION TO CUDA C++


- ▶ What will you learn in this session?
  - ▶ Start with vector addition
  - ▶ Write and launch CUDA C++ kernels
  - ▶ Manage GPU memory
  - ▶ (Manage communication and synchronization)-> next session
- ▶ (Some knowledge of C or C++ programming is assumed.)

# HETEROGENEOUS COMPUTING


- ▶ **Host** The CPU and its memory (host memory)
- ▶ **Device** The GPU and its memory (device memory)




# PORTING TO CUDA




# SIMPLE PROCESSING FLOW



# SIMPLE PROCESSING FLOW




# SIMPLE PROCESSING FLOW



# PARALLEL PROGRAMMING IN CUDA C++

- ▶ GPU computing is about massive parallelism!
- ▶ We need an interesting example...
- ▶ We'll start with vector addition



# GPU KERNELS: DEVICE CODE

```
__global__ void mykernel(void) {  
}
```

- ▶ CUDA C++ keyword **\_\_global\_\_** indicates a function that:
  - ▶ Runs on the device
  - ▶ Is called from host code (can also be called from other device code)
- ▶ **nvcc** separates source code into host and device components
  - ▶ Device functions (e.g. **mykernel ()**) processed by NVIDIA compiler
  - ▶ Host functions (e.g. **main ()**) processed by standard host compiler:
    - ▶ **gcc, cl.exe**

# GPU KERNELS: DEVICE CODE

```
mykernel<<<1,1>>>();
```

- ▶ Triple angle brackets mark a call to *device* code
  - ▶ Also called a “kernel launch”
  - ▶ We’ll return to the parameters (1,1) in a moment
  - ▶ The parameters inside the triple angle brackets are the CUDA kernel **execution configuration**
- ▶ That’s all that is required to execute a function on the GPU!

# MEMORY MANAGEMENT

- ▶ Host and device memory are separate entities

- ▶ *Device* pointers point to GPU memory

Typically passed to device code

Typically *not* dereferenced in host code

- ▶ *Host* pointers point to CPU memory

Typically not passed to device code


Typically *not* dereferenced in device code

- ▶ (Special cases: Pinned pointers, ATS, managed memory)

- ▶ Simple CUDA API for handling device memory

- ▶ `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`

- ▶ Similar to the C equivalents `malloc()`, `free()`, `memcpy()`



# RUNNING CODE IN PARALLEL

- ▶ GPU computing is about massive parallelism
  - ▶ So how do we run code in parallel on the device?

```
add<<< 1, 1 >>>();
```



```
add<<< N, 1 >>>();
```

- ▶ Instead of executing `add()` once, execute  $N$  times in parallel

# VECTOR ADDITION ON THE DEVICE

- ▶ With **add()** running in parallel we can do vector addition
- ▶ Terminology: each parallel invocation of **add()** is referred to as a **block**
  - ▶ The set of all blocks is referred to as a **grid**
  - ▶ Each invocation can refer to its block index using **blockIdx.x**

```
__global__ void add(int *a, int *b, int *c) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

- ▶ By using **blockIdx.x** to index into the array, each block handles a different index
- ▶ Built-in variables like **blockIdx.x** are zero-indexed (C/C++ style), 0..**N**-1, where **N** is from the kernel execution configuration indicated at the kernel launch

# VECTOR ADDITION ON THE DEVICE

```
#define N 512

int main(void) {
    int *a, *b, *c;                      // host copies of a, b, c
    int *d_a, *d_b, *d_c;                // device copies of a, b, c
    int size = N * sizeof(int);

    // Alloc space for device copies of a, b, c
    cudaMalloc((void **) &d_a, size);
    cudaMalloc((void **) &d_b, size);
    cudaMalloc((void **) &d_c, size);

    // Alloc space for host copies of a, b, c and setup input values
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);
```

# VECTOR ADDITION ON THE DEVICE

```
// Copy inputs to device
cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);
// Launch add() kernel on GPU with N blocks
add<<<N,1>>>(d_a, d_b, d_c);

// Copy result back to host
cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);

// Cleanup
free(a); free(b); free(c);
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
return 0;
}
```

# REVIEW (1 OF 2)

- ▶ Difference between *host* and *device*
  - ▶ *Host* CPU
  - ▶ *Device* GPU
- ▶ Using **\_\_global\_\_** to declare a function as device code
  - ▶ Executes on the device
  - ▶ Called from the host (or possibly from other device code)
- ▶ Passing parameters from host code to a device function

# REVIEW (2 OF 2)

- ▶ Basic device memory management
  - ▶ `cudaMalloc()`
  - ▶ `cudaMemcpy()`
  - ▶ `cudaFree()`
- ▶ Launching parallel kernels
  - ▶ Launch **N** copies of `add()` with `add<<<N,1>>>(...);`
  - ▶ Use `blockIdx.x` to access block index

# CUDA THREADS

- ▶ Terminology: a block can be split into parallel **threads**
- ▶ Let's change `add()` to use parallel *threads* instead of parallel *blocks*

```
__global__ void add(int *a, int *b, int *c) {  
    c[threadIdx.x] = a[threadIdx.x] + b[threadIdx.x];  
}
```

- ▶ We use **threadIdx.x** instead of **blockIdx.x**
- ▶ Need to make one change in **main()**:


```
add<<< 1, N >>>();
```

# COMBINING BLOCKS AND THREADS

- ▶ We've seen parallel vector addition using:
  - ▶ Many blocks with one thread each
  - ▶ One block with many threads
- ▶ Let's adapt vector addition to use both *blocks* and *threads*
- ▶ Why? We'll come to that...
- ▶ First let's discuss data indexing...

# INDEXING ARRAYS WITH BLOCKS AND THREADS

- ▶ No longer as simple as using `blockIdx.x` and `threadIdx.x`
  - ▶ Consider indexing an array with one element per thread (8 threads/block):




- ▶ With  $M$  threads/block a unique index for each thread is given by:

```
int index = threadIdx.x + blockIdx.x * M;
```

# INDEXING ARRAYS: EXAMPLE

- ▶ Which thread will operate on the red element?



```
int index = threadIdx.x + blockIdx.x * M;  
= 5 + 2 * 8;  
= 21;
```

# VECTOR ADDITION WITH BLOCKS AND THREADS

- ▶ Use the built-in variable `blockDim.x` for threads per block

```
int index = threadIdx.x + blockIdx.x * blockDim.x
```

- ▶ Combined version of `add()` to use parallel threads *and* parallel blocks:

```
__global__ void add(int *a, int *b, int *c) {
    int index = threadIdx.x + blockIdx.x * blockDim.x;
    c[index] = a[index] + b[index];
}
```

- ▶ What changes need to be made in `main()`?

# ADDITION WITH BLOCKS AND THREADS

```
#define N (2048*2048)
#define THREADS_PER_BLOCK 512
int main(void) {
    int *a, *b, *c;                                // host copies of a, b, c
    int *d_a, *d_b, *d_c;                            // device copies of a, b, c
    int size = N * sizeof(int);
    // Alloc space for device copies of a, b, c
    cudaMalloc((void **) &d_a, size);
    cudaMalloc((void **) &d_b, size);
    cudaMalloc((void **) &d_c, size);
    // Alloc space for host copies of a, b, c and setup input values
    a = (int *)malloc(size); random_ints(a, N);
    b = (int *)malloc(size); random_ints(b, N);
    c = (int *)malloc(size);
```

# ADDITION WITH BLOCKS AND THREADS

```
// Copy inputs to device
cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

// Launch add() kernel on GPU
add<<<N/THREADS_PER_BLOCK,THREADS_PER_BLOCK>>>(d_a, d_b, d_c);

// Copy result back to host
cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);

// Cleanup
free(a); free(b); free(c);
cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
return 0;
}
```

# HANDLING ARBITRARY VECTOR SIZES

- ▶ Typical problems are not friendly multiples of `blockDim.x`
- ▶ Avoid accessing beyond the end of the arrays:

```
__global__ void add(int *a, int *b, int *c, int n) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    if (index < n)  
        c[index] = a[index] + b[index];  
}
```

- ▶ Update the kernel launch:

```
add<<< (N + M-1) / M>>>(d_a, d_b, d_c, N);
```

# WHY BOTHER WITH THREADS?

- ▶ Threads seem unnecessary
  - ▶ They add a level of complexity
  - ▶ What do we gain?
- ▶ Unlike parallel blocks, threads have mechanisms to:
  - ▶ Communicate
  - ▶ Synchronize
- ▶ To look closer, we need a new example... (next session)

# REVIEW

- ▶ Launching parallel kernels
  - ▶ Launch N copies of add ( ) with add<<<N/M, M>>> ( ... ) ;
  - ▶ Use **blockIdx.x** to access block index
  - ▶ Use **threadIdx.x** to access thread index within block
- ▶ Assign elements to threads:

```
int index = threadIdx.x + blockIdx.x * blockDim.x;
```

# FUTURE SESSIONS

- ▶ CUDA Shared Memory
- ▶ CUDA GPU architecture and basic optimizations
- ▶ Atomics, Reductions, Warp Shuffle
- ▶ Using Managed Memory
- ▶ Concurrency (streams, copy/compute overlap, multi-GPU)
- ▶ Analysis Driven Optimization
- ▶ Cooperative Groups

# FURTHER STUDY

- ▶ An introduction to CUDA:
  - ▶ <https://devblogs.nvidia.com/easy-introduction-cuda-c-and-c/>
- ▶ Another introduction to CUDA:
  - ▶ <https://devblogs.nvidia.com/even-easier-introduction-cuda/>
- ▶ CUDA Programming Guide:
  - ▶ <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>
- ▶ CUDA Documentation:
  - ▶ <https://docs.nvidia.com/cuda/index.html>
  - ▶ <https://docs.nvidia.com/cuda/cuda-runtime-api/index.html> (runtime API)

# HOMEWORK

- ▶ Log into Summit (ssh [username@home.ccs.ornl.gov](mailto:username@home.ccs.ornl.gov) -> ssh summit)
- ▶ Clone GitHub repository:
  - ▶ Git clone [git@github.com:olcf/cuda-training-series.git](https://github.com/olcf/cuda-training-series.git)
- ▶ Follow the instructions in the readme.md file:
  - ▶ <https://github.com/olcf/cuda-training-series/blob/master/exercises/hw1/readme.md>
- ▶ Prerequisites: basic linux skills, e.g. ls, cd, etc., knowledge of a text editor like vi/emacs, and some knowledge of C/C++ programming



# QUESTIONS?

